

Project 2

Importing the data and saving it

```
bank <- read.csv(file = '/Users/trishalvarma/Desktop/bank-full.csv')
```

Data Cleaning

Website: <https://archive.ics.uci.edu/ml/datasets/bank+marketing> From: "UCI Machine Learning Repository" — Used any(is.na(data)) to see if there are any data missing from the columns. Then used na.omit, because data can be filled with NA and not show up as a blank data. Eliminates missing values. Provided the link from where the data is downloaded.

Steps taken are listed above, to remove any blank/na values. #Data Cleanup

```
bank$default <- NULL
bank$day <- NULL

bank$duration <- NULL
bank$pdays <- NULL
bank$previous <- NULL
bank$poutcome <- NULL

bank$y <- NULL
bank$contact <- NULL
str(bank)

## 'data.frame': 45211 obs. of 9 variables:
## $ age : int 58 44 33 47 33 35 28 42 58 43 ...
## $ job : chr "management" "technician" "entrepreneur" "blue-collar"
## ...
## $ marital : chr "married" "single" "married" "married" ...
## $ education: chr "tertiary" "secondary" "secondary" "unknown" ...
## $ balance : int 2143 29 2 1506 1 231 447 2 121 593 ...
## $ housing : chr "yes" "yes" "yes" "yes" ...
## $ loan : chr "no" "no" "yes" "no" ...
## $ month : chr "may" "may" "may" "may" ...
## $ campaign : int 1 1 1 1 1 1 1 1 1 1 ...
```

Grouping bank and bank job so we can use it as combined.

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

group_by(bank, bank$job)

## # A tibble: 45,211 x 10
## # Groups:   bank$job [12]
##   age job   marital education balance housing loan  month campaign
##   <int> <chr> <chr>   <chr>      <int> <chr>  <chr> <chr>   <int>
##   <chr>
## 1    58 mana... married tertiary    2143 yes    no    may        1
management
## 2    44 tech... single  secondary     29 yes    no    may        1
technician
## 3    33 entr... married secondary     2 yes    yes   may        1
entreprene...
## 4    47 blue... married unknown    1506 yes    no    may        1
blue-coll...
## 5    33 unkn... single  unknown      1 no     no    may        1
unknown
## 6    35 mana... married tertiary    231 yes    no    may        1
management
## 7    28 mana... single  tertiary    447 yes    yes   may        1
management
## 8    42 entr... divorc... tertiary     2 yes    no    may        1
entreprene...
## 9    58 reti... married primary    121 yes    no    may        1
retired
## 10   43 tech... single  secondary    593 yes    no    may        1
technician
## # ... with 45,201 more rows

bank$campaign <- NULL
```

Creating factors of the variables we will be using.

We also removed the outliers from the bank dataset.

```
bank$education <- as.factor(bank$education)
bank$marital <- as.factor(bank$marital)
bank$job <- as.factor(bank$job)
bank$housing <- as.factor(bank$housing)
bank$loan <- as.factor(bank$loan)

outliers <- boxplot(bank$age, plot=FALSE)$out
```

```
bank <- bank[-which(bank$age %in% outliers),]
nrow(bank)

## [1] 44724

outliers1 <- boxplot(bank$balance, plot=FALSE)$out
bank <- bank[-which(bank$balance %in% outliers1),]
nrow(bank)

## [1] 40028
```

Data Exploration

Using more than 5 functions to look at the Air Quality data, then providing a informative R graphs using a histogram, and a plot chart to show the just the sheer amount of data that was collected.

```
tapply(bank$age, bank$job, mean)

##      admin.  blue-collar  entrepreneur  housemaid  management
## 38.95821    39.82201    41.78245    45.95375    40.12708
##  retired self-employed  services  student  technician
## 57.76081    40.02168    38.54024    26.46659    39.17989
## unemployed  unknown
## 40.68031    46.92946

# then we also check on balance, and education
tapply(bank$balance, bank$education, mean)

## primary secondary tertiary unknown
## 601.0054  589.1953  713.5273  698.7475

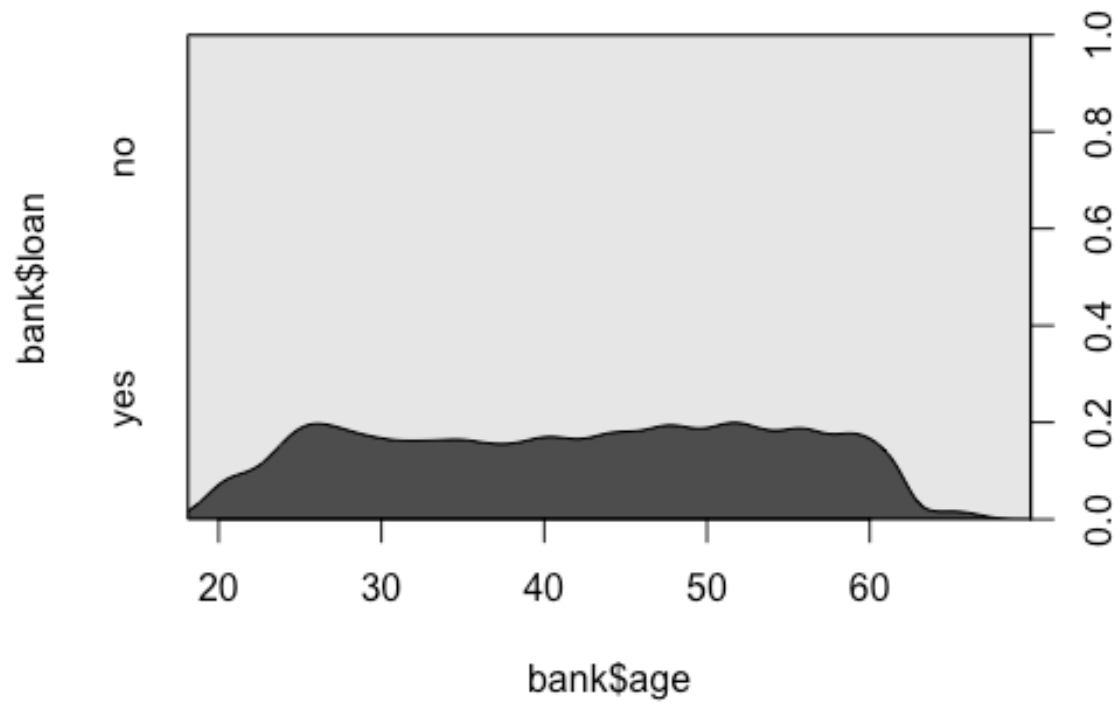
# we can create a table and combine the marital status and the job they have.
library(gmodels)
bn.table <- table(bank$marital, bank$job)
bn.table

##
##      admin. blue-collar  entrepreneur  housemaid  management  retired
## divorced    706        702        166        154        966        287
## married    2412       6395        944        800       4591       1197
## single     1596       1836        223        127       2556        88
##
##      self-employed  services  student  technician  unemployed  unknown
## divorced         113        515         6        865        154        13
## married          873       2200         49       3602        640       168
## single          398       1124       783       2365        354        60
```


Now we can see visual data exploration.

Here we are able to see the loan and age taken, and the loan and balance taken.

```
cdplot(bank$loan~bank$age)
```

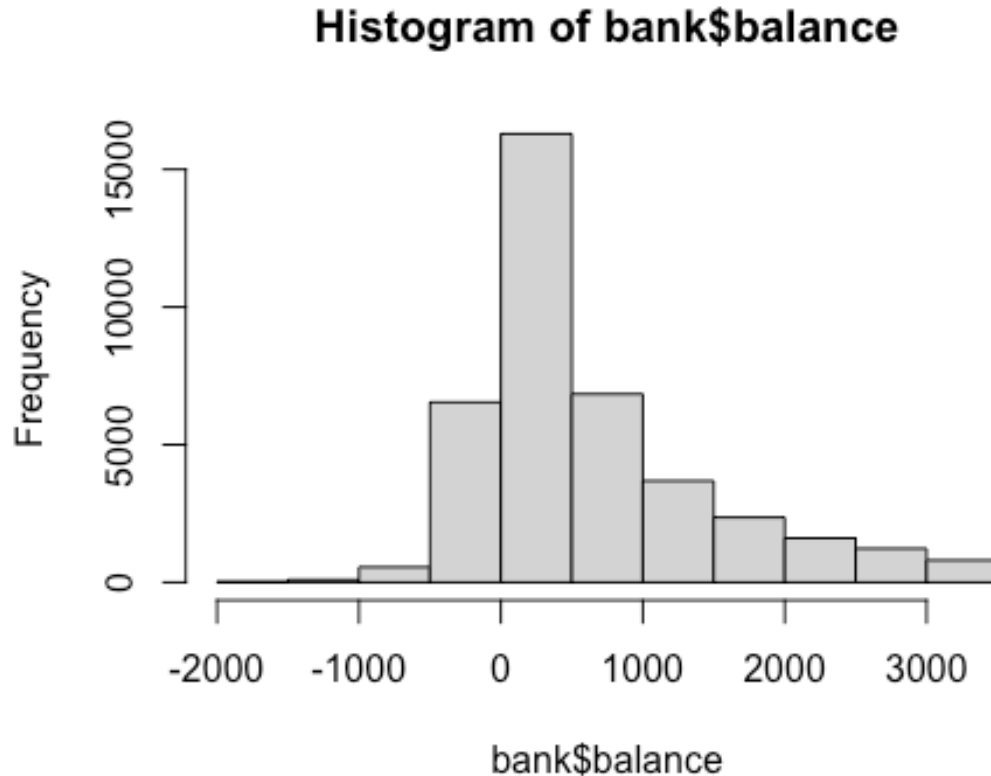


```
cdplot(bank$loan~ bank$balance)
```



second exploration visually is the histogram

`hist(bank$balance)` *#Here we notice that there are a few that shows negative balance, and could effect the algorithms we will be running.*



Lets

create a test and a training data set we create 2 training and 2 test, just in case.

```
set.seed(1234)
bank1 <- bank
bank1 <- na.omit(bank1)

sample_i <- sample(1:nrow(bank1), .80*nrow(bank1), replace=FALSE)
train <- bank1[sample_i,]
test <- bank1[-sample_i,]

train1 <- bank1[sample_i,]
test1 <- bank1[-sample_i,]
```

For classification algorithm, we will run a Logistical regression. Even though it is a regression, it is actually a clasfication algorithm.


```
gg <- glm(age~., data = train)
summary(gg)


##
## Call:
## glm(formula = age ~ ., data = train)
##
## Deviance Residuals:
```


```


##      Min      1Q      Median      3Q      Max
## -25.0823  -5.9086  -0.8489    5.3417   27.9973
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    4.680e+01  2.941e-01 159.139 < 2e-16 ***
## jobblue-collar -9.911e-01  1.693e-01  -5.855 4.82e-09 ***
## jobentrepreneur 1.508e+00  2.846e-01   5.299 1.17e-07 ***
## jobhousemaid    2.754e+00  3.107e-01   8.865 < 2e-16 ***
## jobmanagement   1.422e+00  1.902e-01   7.474 8.01e-14 ***
## jobretired      1.460e+01  2.664e-01  54.799 < 2e-16 ***
## jobself-employed 5.632e-01  2.808e-01   2.005  0.0449 *
## jobservices     -9.835e-01  1.947e-01  -5.052 4.39e-07 ***
## jobstudent      -9.244e+00  3.473e-01 -26.620 < 2e-16 ***
## jobtechnician    7.339e-02  1.714e-01   0.428  0.6686
## jobunemployed    3.260e-01  2.966e-01   1.099  0.2717
## jobunknown       3.866e+00  5.852e-01   6.607 4.00e-11 ***
## maritalmarried  -2.138e+00  1.445e-01 -14.794 < 2e-16 ***
## maritalsingle   -9.605e+00  1.581e-01 -60.747 < 2e-16 ***
## educationsecondary -3.024e+00  1.432e-01 -21.120 < 2e-16 ***
## educationtertiary -4.365e+00  1.802e-01 -24.223 < 2e-16 ***
## educationunknown  3.319e-01  2.597e-01   1.278  0.2012
## balance          6.935e-04  5.485e-05  12.643 < 2e-16 ***
## housingyes      -2.103e+00  1.067e-01 -19.705 < 2e-16 ***
## loanyes         -1.168e-01  1.215e-01  -0.961  0.3365
## monthaug        1.347e+00  2.236e-01   6.025 1.71e-09 ***
## monthdec        1.592e-01  7.241e-01   0.220  0.8259
## monthfeb        3.761e-01  2.572e-01   1.463  0.1436
## monthjan        4.482e-01  3.064e-01   1.462  0.1436
## monthjul        2.936e-01  2.119e-01   1.386  0.1658
## monthjun        1.115e+00  2.224e-01   5.016 5.31e-07 ***
## monthmar        2.088e-02  5.029e-01   0.042  0.9669
## monthmay        1.706e-01  1.940e-01   0.880  0.3791
## monthnov        1.352e+00  2.413e-01   5.604 2.12e-08 ***
## monthoct        1.704e+00  4.230e-01   4.028 5.63e-05 ***
## monthsep       -1.072e-01  4.640e-01  -0.231  0.8173
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 63.64874)
##
##      Null deviance: 3152156  on 32021  degrees of freedom
## Residual deviance: 2036187  on 31991  degrees of freedom
## AIC: 223907
##
## Number of Fisher Scoring iterations: 2
plot(gg)

```









```
pred1 <- predict(gg, newdata = test)

cor1 <- cor(pred1, test$age)
mse1 <- mean((pred1 - test$age)^2)

print(paste("MSE: ", mse1))
## [1] "MSE: 61.6497266522616"

print(paste("Corrleation: ", cor1))
## [1] "Corrleation: 0.595206875117885"
```

Not bad, but not good, we have a MSE of 61, and a correlation of 59.52 but almost 60. So we'll keep that in mind for later.

#Let's run our second algorithm kNN, and this time using test1 and train1.

```
test1$housing <- as.numeric(test1$housing)
train1$housing <- as.numeric(train1$housing)

test1$age <- as.numeric(test1$age)
train1$age <- as.numeric(train1$age)
```

```

test1$job <- as.numeric(test1$job)
train1$job <- as.numeric(train1$job)

test1$marital <- as.numeric(test1$marital)
train1$marital <- as.numeric(train1$marital)

test1$education <- as.numeric(test1$education)
train1$education <- as.numeric(train1$education)

test1$balance <- as.numeric(test1$balance)
train1$balance <- as.numeric(train1$balance)

test1$loan <- as.numeric(test1$loan)
train1$loan <- as.numeric(train1$loan)

library(caret)

## Loading required package: lattice

## Loading required package: ggplot2

library(lattice)
library(ggplot2)

knn_fit <- knnreg(age~., data = test1, k = 5)
pred_knn <- predict(knn_fit, test1)

cor_knn <- cor(pred_knn, test1$age)
mse_knn <- mean((pred_knn - test1$age)^2)

print(paste("Correlation for KNN: ", knn_fit))

## [1] "Correlation for KNN:  list(y = c(`3` = 33, `5` = 33, `20` = 33, `21`
= 28, `22` = 56, `28` = 52, `29` = 46, `30` = 36, `37` = 25, `40` = 37, `41`
= 44, `49` = 55, `59` = 40, `63` = 57, `68` = 59, `72` = 31, `75` = 43, `78`
= 55, `91` = 42, `94` = 60, `96` = 36, `98` = 60, `102` = 53, `103` = 52,
`104` = 59, `109` = 59, `112` = 46, `116` = 44, `117` = 41, `118` = 33, `120`
= 57, `122` = 51, `126` = 33, `131` = 55, `132` = 32, `133` = 38, `141` = 53,
`149` = 43, `151` = 51, `158` = 60, `159` = 52, `176` = 53, `178` = 34,
\n`185` = 36, `191` = 51, `197` = 38, `203` = 44, `213` = 59, `220` = 39,
`226` = 48, `229` = 36, `236` = 45, `248` = 40, `249` = 39, `252` = 53, `258`
= 30, `270` = 40, `271` = 42, `274` = 56, `279` = 38, `281` = 50, `289` = 32,
`290` = 40, `295` = 34, `302` = 51, `304` = 36, `313` = 55, `317` = 36, `318`
= 51, `324` = 32, `333` = 30, `337` = 42, `341` = 41, `347` = 45, `350` = 55,
`354` = 36, `361` = 48, `362` = 42, `363` = 27, `372` = 38, `374` = 25, `379`
= 33, `390` = 58, `394` = 27, `400` = 47, \n`402` = 48, `405` = 52, `411` =
32, `414` = 48, `421` = 49, `428` = 54, `432` = 42, `435` = 32, `446` = 51,
`451` = 50, `460` = 35, `462` = 39, `477` = 50, `481` = 51, `482` = 41, `486`
= 54, `489` = 47, `495` = 30, `504` = 29, `507` = 57, `512` = 34, `518` = 44,
= 27, `8696` = 34, `8698` = 26, `8718` = 29, `8722` = 38, `8728` = 34, `8729`

```

```

= 42, `8738` = 54, `8742` = 41, `8746` = 33, `8752` = 59, `8759` = 28, `8760`
= 29, `8767` = 29, `8780` = 34, `8783` = 29, `8784` = 26, `8789` = 22, `8801`
= 21, `8809` = 26, `8811` = 24, `8816` = 52, `8822` = 29, `8823` = 29, `8834`
= 32, `8835` = 36, `8836` = 47, `8839` = 40, `8843` = 57, `8847` = 36, `8848`
= 45, `8849` = 30, `8862` = 35, `8876` = 39, `8879` = 53, `8891` = 46, `8894`
## [2] "Correlation for KNN: 5"
## [3] "Correlation for KNN: age ~ job + marital + education + balance +
housing + loan + month"
## [4] "Correlation for KNN: list()"
## [5] "Correlation for KNN: list()"

print(paste("MSE for KNN: ", mse_knn))

## [1] "MSE for KNN: 67.6595729200846"

```

This prints out a long data, but we can see that the mse is 67, and the correlation is 5. kNN might not be the best suited for this data.

Lets run a random forest and see our tree

```

library(randomForest)

## randomForest 4.6-14

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:ggplot2':
##
##     margin

## The following object is masked from 'package:dplyr':
##
##     combine

set.seed(1234)
tree <- randomForest(age~., data=test, importance = TRUE)
tree

##
## Call:
## randomForest(formula = age ~ ., data = test, importance = TRUE)
##              Type of random forest: regression
##              Number of trees: 500
## No. of variables tried at each split: 2
##
##              Mean of squared residuals: 60.52472
##              % Var explained: 36.6

```

#Analysis

Ranking of algorithms

- 1) random Forest
- 2) Linear Regression
- 3) kNN

In this case as well, we have random Forest being top, because of the number of trees it can create. We can increase the number of tries, but 1 is sufficient to know that it performed well. The mean residuals. Linear regression also gave us very good number, as commented above, and kNN did give us a better MSE, however it was very messy to deal with. #the data for kNN was edited so it could fit in the pdf and not be long. The time it took for the kNN as well on such a large data also effected it performance.

We were able to look at the age and loan amount for people and in this case the random forrest would be the best to analyze the data.

What we learned; EX: If people older age have higher balance, then it is likely they would be approved for loans, whereas, a lower balance made it a risk for loans. It was definitely a good data to work with, the more we learn, the more we can work with.